

# Coursework 3

36071280

02/11/2021

```
library(tidyverse)
```

## Question 1

Q1. Read in the data and print the dimensions of the Storm Events data frame

```
storm_events <- read_csv("Australia_severe_storms_1975-2015.csv", show_col_types = FALSE)
```

```
## New names:
## * ` ` -> ...10
## * ` ` -> ...11
## * ` ` -> ...12
## * ` ` -> ...13
## * ` ` -> ...14
```

```
print(dim(storm_events))
```

```
## [1] 14457    14
```

## Question 2

Q2. Clean the data by removing the variable ID and also Waterspout events from the database. Print the dimensions of the cleaned data frame. Also print the first few rows without the 6 columns of comments, without creating an intermediate data frame.

```
# Removing variable ID and Waterspout event
storm_events_clean <- storm_events %>% select(-(ID)) %>% filter(storm_events$Database != "Waterspout")
print(dim(storm_events_clean))
```

```
## [1] 14417    13
```

#Q3. Add a column to your data frame containing the time zone of each event using the following OlsonNames() classifications.

```
#declaring variables for timezone mapping

list_of_relevant_australian_tz <- c("QLD" = "Australia/Queensland",
                                     "NSW" = "Australia/NSW",
                                     "NSW_BrokenHill" = "Australia/Broken_Hill",
                                     "VIC" = "Australia/Victoria",
                                     "SA" = "Australia/South",
                                     "WA" = "Australia/West",
                                     "TAS" = "Australia/Tasmania",
                                     "NT" = "Australia/North",
```

```

      "ACT" = "Australia/ACT")

new_south_wales = "NSW"
aus_central_time = "ACT"
broken_hill = "NSW_BrokenHill"
broken_hill_expr = "broken"

allocate_tz <-function(state, nearest_town) {

  if (length(state) && !is.na(state))
  {
    if (state != new_south_wales)
      return (list_of_relevant_australian_tz[state])
    else if (state == new_south_wales)
    {

      if (length(nearest_town) && !is.na(nearest_town) && str_detect(tolower(nearest_town), broken_hill_expr))
        return (list_of_relevant_australian_tz[broken_hill])
      else
        return (list_of_relevant_australian_tz[new_south_wales])
    }
  }
  return (list_of_relevant_australian_tz[aus_central_time])
}

storm_events_tz <- storm_events_clean %>% mutate(AustralianTimeZone = NA)

for ( i in 1:nrow(storm_events_clean))
{
  storm_events_tz$AustralianTimeZone[i] <- allocate_tz(storm_events_clean$State[i], storm_events_clean$
}

#storm_events_tz2 = lapply(storm_events_tz, 2, FUN = allocate_tz)

```

#### Q4. Parse the date, time and time zones from the necessary columns to create a new variable in the data

frame which converts the time into UTC. You may need the function `lubridate::as_datetime()` and/or the use of loops. Print the first few rows of the resultant data frame, without the 6 columns of comments, again without creating an intermediate data frame.

```

library(lubridate)

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

storm_events_utc = data.frame()
#storm_events_utc <- storm_events_tz %>% mutate(UTCDateTime = NA)

for ( s in list_of_relevant_australian_tz)

```

```
{
  #print(s)
  storm_events_utc <- rbind(storm_events_utc, storm_events_tz %>% filter(storm_events_tz$AustralianTime
    #storm_events_utc$UTCDateTime
}
storm_events_utc %>% arrange(storm_events_utc$`Event ID`)
```

```
## # A tibble: 14,417 x 15
##   `Event ID` Database `Date/Time`   `Nearest town`   State Latitude Longitude
##   <dbl> <chr>    <chr>          <chr>          <chr>    <dbl>    <dbl>
## 1    20812 Wind     23/11/1975 0~ SYDNEY      NSW      -33.9     151.
## 2    20813 Tornado 02/12/1975 1~ BARHAM      NSW      -35.6     144.
## 3    20814 Wind     09/01/1976 0~ COFF'S HARBOUR NSW      -30.3     153.
## 4    20815 Hail     16/02/1976 1~ BANKSTOWN    NSW      -33.9     151.
## 5    20816 Rain     25/10/1976 1~ BOOMI        NSW      -28.4     153.
## 6    20817 Hail     08/11/1976 1~ YOUNG        NSW      -34.3     148.
## 7    20818 Hail     09/11/1976 1~ SYDNEY (WESTERN S~ NSW      -33.7     151
## 8    20818 Rain     09/11/1976 1~ WESTERN SUBURBS NSW      -33.7     151
## 9    20819 Wind     09/11/1976 1~ SYDNEY      NSW      -33.9     151.
## 10   20820 Wind     10/11/1976 1~ MOREE/TRANGIE NSW      -29.5     150.
## # ... with 14,407 more rows, and 8 more variables: Comments <chr>, ...10 <chr>,
## #   ...11 <chr>, ...12 <lg1>, ...13 <lg1>, ...14 <lg1>,
## #   AustralianTimeZone <chr>, UTCDateTime <dtm>
```

```
storm_events_utc2 <- storm_events_tz %>% mutate(UTCDateTime = NA)

for ( j in 1: nrow(storm_events_utc2)) {
  storm_events_utc2$UTCDateTime[j] <- (dmy_hm(storm_events_utc2$`Date/Time`[j], tz=storm_events_utc2$Au
}
storm_events_utc2$UTCDateTime <- storm_events_utc2$UTCDateTime %>% as_datetime()

#comment_columns = c(...10,...11, ...12,...13, ...14, Comments)
# filter all the tz by the name,
# mass apply dmy_hm -> loop should be 1:10
head(storm_events_utc2 %>% select(-c(...10,...11, ...12,...13, ...14, Comments)))
```

```
## # A tibble: 6 x 9
##   `Event ID` Database `Date/Time`   `Nearest town`   State Latitude Longitude
##   <dbl> <chr>    <chr>          <chr>          <chr>    <dbl>    <dbl>
## 1    20812 Wind     23/11/1975 07:00 SYDNEY      NSW      -33.9     151.
## 2    20813 Tornado 02/12/1975 14:00 BARHAM      NSW      -35.6     144.
## 3    20814 Wind     09/01/1976 08:50 COFF'S HARBOUR NSW      -30.3     153.
## 4    20815 Hail     16/02/1976 14:00 BANKSTOWN    NSW      -33.9     151.
## 5    20816 Rain     25/10/1976 14:00 BOOMI        NSW      -28.4     153.
## 6    20817 Hail     08/11/1976 14:00 YOUNG        NSW      -34.3     148.
## # ... with 2 more variables: AustralianTimeZone <chr>, UTCDateTime <dtm>

storm_events_utc = storm_events_utc2
```

#Q5. 5. Create new variables for the month and year of each event. Print the first few rows of the resultant data frame, without the 6 columns of comments and without creating an intermediate data frame.

```
storm_events_with_cols = storm_events_utc %>% mutate(month_storm = month(storm_events_utc$UTCDateTime),
  year_storm = year(storm_events_utc$UTCDateTime))
```

```
#...10,...11, ...12,...13, ...14, Comments
print(storm_events_with_cols[-c(8:13)])
```

```
## # A tibble: 14,417 x 11
##   `Event ID` Database `Date/Time`   `Nearest town`   State Latitude Longitude
##   <dbl> <chr>    <chr>         <chr>         <chr>    <dbl>    <dbl>
## 1    20812 Wind      23/11/1975 0~ SYDNEY      NSW      -33.9    151.
## 2    20813 Tornado  02/12/1975 1~ BARHAM      NSW      -35.6    144.
## 3    20814 Wind      09/01/1976 0~ COFF'S HARBOUR NSW      -30.3    153.
## 4    20815 Hail      16/02/1976 1~ BANKSTOWN    NSW      -33.9    151.
## 5    20816 Rain      25/10/1976 1~ BOOMI        NSW      -28.4    153.
## 6    20817 Hail      08/11/1976 1~ YOUNG         NSW      -34.3    148.
## 7    20818 Hail      09/11/1976 1~ SYDNEY (WESTERN S~ NSW      -33.7    151
## 8    20818 Rain      09/11/1976 1~ WESTERN SUBURBS NSW      -33.7    151
## 9    20819 Wind      09/11/1976 1~ SYDNEY        NSW      -33.9    151.
## 10   20820 Wind      10/11/1976 1~ MOREE/TRANGIE NSW      -29.5    150.
## # ... with 14,407 more rows, and 4 more variables: AustralianTimeZone <chr>,
## #   UTCDateTime <dtm>, month_storm <dbl>, year_storm <dbl>
```

**Q6. After discarding Waterspout events there are five types of events left in the data; Rain, Hail, Lighting, Wind, and Tornado.**


##i) Create a new data frame which contains the total number of counts for each of the above type of events for each of the twelve months over the forty year period.

```
storm_cols_grouped = storm_events_with_cols %>% group_by(Database, month_storm)
storm_counts_df = data.frame( count(storm_cols_grouped))
```

ii) On a single plot, plot the total number of counts of each event against month. Use the R object month.abb for the labels of the months in the plot.

```
storm_counts_df$month_storm <- factor(storm_counts_df$month_storm, labels=month.abb)

ggplot(storm_counts_df,aes(factor(Database), n, fill=month_storm )) +
  geom_bar(position="dodge2", stat="identity")
```



```
#ggplot(storm_counts_df,aes(factor(Database), n, fill=month_storm )) +
# geom_line()
```

Q7. From the answer to Question 5, the 6 columns titled Comments, X, X.1, X.2, X.3, X.4 consist of comments.

i) Combine the comments from these columns into a single column, named All.comments.

```
storm_events_with_cols[is.na(storm_events_with_cols)] <- ''
storm_counts_all_comments <- storm_events_with_cols %>% mutate(AllComments = str_trim(paste(storm_event,
                                                                                             storm_events_with_col,
                                                                                             storm_events_with_col)))
#glimpse(storm_counts_all_comments)
```

ii) Select the following columns to keep for further analysis, Event.ID, Database, State, All.comments, and the year variable you created.

```
#storm_counts_all_comments$`Event ID`
#select(storm_counts_all_comments, storm_counts_all_comments$`Event ID`)
storm_analysis <- storm_counts_all_comments %>% select(c(`Event ID`, Database, State, AllComments, year))
print(sapply(storm_analysis, class))
```

```
##      Event ID      Database      State AllComments  year_storm
##      "numeric" "character" "character" "character"  "numeric"
```

**Q8.** Now we use the answer to Question 7(ii) for further analysis

i) Create an indicator variable which states whether or not a storm event has resulted in a flash flood. Make sure you sort out all terms relating to flash floods.

```
flash_flood_expr <- "flash.*?flood*" # flash - any? - flood
flash_flood_expr_space_mul <- "flash\\s*?flood*" # flash - space(s)? - flood
flash_flood_expr_mul <- "flash(.|\\s*)?flood*" #final

storm_analysis <- storm_analysis %>% mutate(flash_flood_present = str_detect(storm_analysis$AllComments,
flash_flood_expr_mul))

storm_analysis_with_floods <- storm_analysis %>% filter(flash_flood_present==TRUE)
storm_analysis_grouped <- storm_analysis_with_floods %>% group_by(year_storm)

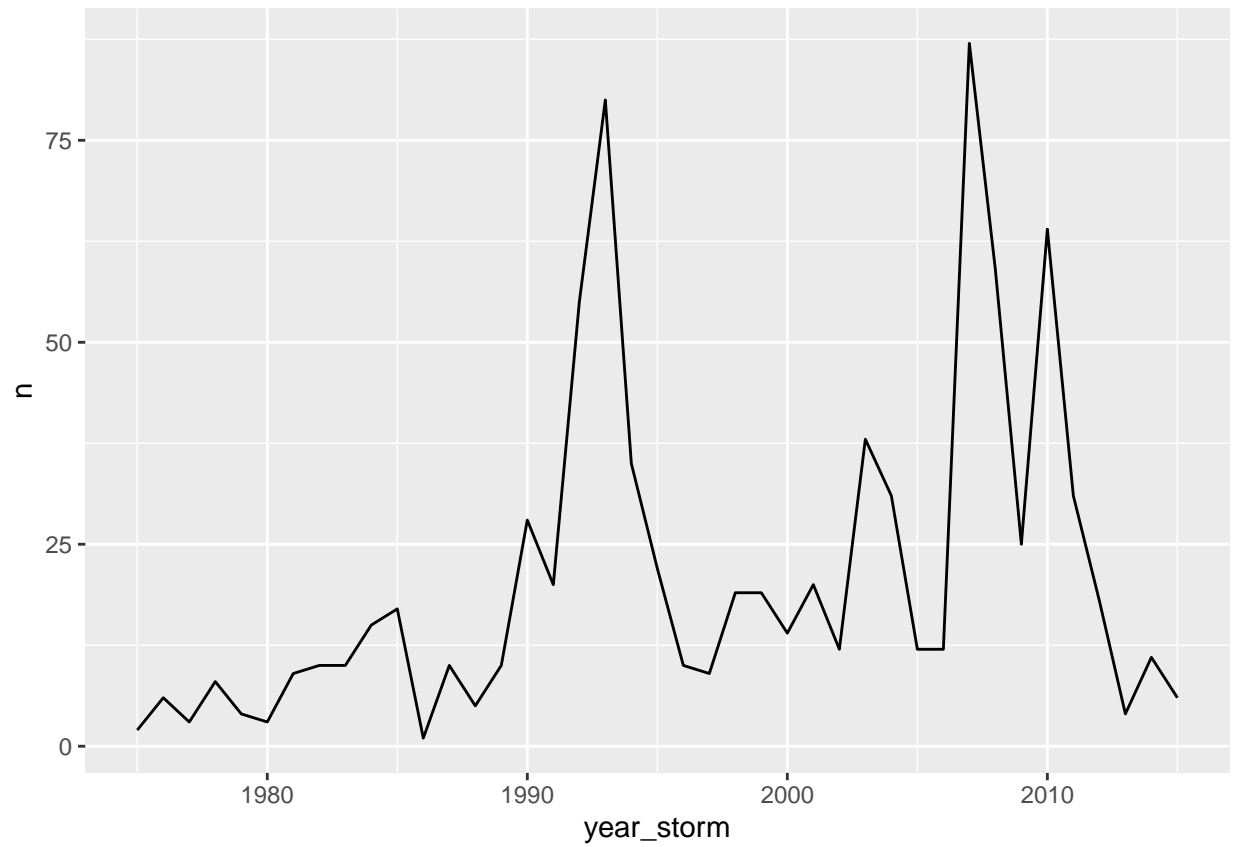
dim(storm_analysis_grouped)

## [1] 854    6
```


ii) Print a plot of the number of flash floods per year from 1975-2015. You may need to first create a vector or data frame to contain the number of flash floods per year.

```
flash_flood_counts <- count(storm_analysis_grouped)

ggplot(flash_flood_counts) + geom_line(aes(year_storm, n))
```



```
ggplot(flash_flood_counts,aes(year_storm, n)) + geom_bar(stat="identity")
```



```
#View(flash_flood_counts)
```

**Q9.** For severe wind events often the wind speed is given. The wind speed is given in knots or km/h.

i) Extract all wind speeds both those in knots and km/h. Hint: Knots can be abbreviated by kts or kt. Also note that wind speed can be a single, double or triple digit number

```
wind_speed_expr <- "(\\d{1,3})\\s?(knot|kt(s)?|km\\s?\\|\\s?h(r)?)"

storm_analysis_speeds <- str_extract_all(storm_analysis, regex(wind_speed_expr))

## Warning in stri_extract_all_regex(string, pattern, simplify = simplify, :
## argument is not an atomic vector; coercing

storm_analysis_speeds <- storm_analysis %>% mutate(wind_speed = str_extract( storm_analysis$AllComments

storm_analysis_speeds <- storm_analysis_speeds %>% filter(!is.na(wind_speed))

#write.csv(storm_analysis_speeds, "storm_analysis_speeds.csv")
```



ii) Convert km/h wind speeds to knots (1 knot = 1.852 km/h) rounding the wind/speed to the nearest knot. Hint: It is helpful to work with a reduced data frame which includes only those observations with a wind speed recorded.

```
knot_to_kmh = 1.852
wind_speed_expr_kmh <- "(\\d{1,3})\\s?(km\\s?\\/\\s?h(r)?)"
extract_number_expr <- "\\d+"

get_wind_value <- function (wind_speed) {
  wind_numeric_value = as.numeric(str_extract(wind_speed, extract_number_expr))
  if (str_detect(wind_speed, regex(wind_speed_expr_kmh, ignore_case = TRUE)))
  {
    return (as.integer(wind_numeric_value/knot_to_kmh))
  }
  else
  {
    return (wind_numeric_value)
  }
}

storm_analysis_speed_numbers <- storm_analysis_speeds %>% mutate(wind_values = NA)
storm_analysis_speed_numbers$wind_values <- sapply(storm_analysis_speed_numbers$wind_speed, FUN = get_w

#as.numeric(str_extract(storm_analysis_speeds$wind_speed[27], extract_number_expr))
#storm_analysis_speed_numbers[1:15,]
```

iii) Print a boxplot of the wind speeds recorded per state.

```
ggplot(storm_analysis_speed_numbers, aes(State, wind_values)) + geom_boxplot(position="dodge2")
```

